

CS4321 - Project 4

Chengcheng Ji (cj368), Pei Xu (px29) and Ella Xue (ex32)

Queries:

1. **SELECT * FROM** Sailors, Reserves **WHERE** Sailors.A <= 500 **AND** Sailors.A >= 100;
- Index used Sailor.A Order 15
2. **SELECT * FROM** Sailors, Reserves
WHERE Sailors.A > 100 **AND** Reserves.H >=100 **AND** Sailors.A = Reserves.H;
- Indexes used Sailor.A Order 15, Reserves.H Order 15
3. **SELECT * FROM** Reserves, Boats
WHERE Boats.E >= 100 **AND** Reserves.H >=700 **AND** Reserves.H > Boats.E;
- Indexes used Boats.E Order 10, Reserves.H Order 15.
4. **SELECT * FROM** Reserves, Boats **WHERE** Boats.E < 5 **AND** Reserves.H >=700;
- Indexes used Boats.E Order 10, Reserves.H Order 15

The schema is:

Sailors A B C

Boats D E F

Reserves G H

Indexes information unclustered:

- Boats E 0 10
- Sailors A 0 15
- Reserves H 0 15

Indexes information clustered:

- Boats E 1 10
- Sailors A 1 15
- Reserves H 1 15

Data description:


Each attribute value was chosen uniformly at random in the range of 0 to 1000; 5000 tuples per relation.

*Data in millisecond

	Full Scan	Uncluster	Cluster
Query 1	15253	15373	9533
Query 2	18607	144	110
Query3	11270	89	36
Query4	2309	45	31

CS4321 - Project 4

Chengcheng Ji (cj368), Pei Xu (px29) and Ella Xue (ex32)



Conclusion: Using BPlusTree indexes, scanning for selection queries is significantly faster than full scan under the queries we have tested. The comparison used the same join method and reader but different scanning method: full scan, un-clustered indexes and clustered indexes.